

Julia for Technical Computing

Josh Day

<https://github.com/joshday>

emailjoshday@gmail.com

To run the examples

```
pkgs = [  
    "Distributions",  
    "Plots",  
    "PyPlot",  
    "GR",  
    "RCall",  
    "BenchmarkTools",  
    "Query",  
    "DataTables"  
]  
  
for pkg in pkgs  
    Pkg.add(pkg)  
end
```

Overview

- What is Julia?
- Design, Type System, and Multiple Dispatch
- Julia Packages
- Macros

What is Julia?

What is Julia?


Julia is a high-level, high-performance dynamic programming language for technical computing, with syntax that is familiar to users of other technical computing environments.

- <http://julialang.org>

What is Julia?

- You'll commonly hear two things from the Julia community:
 - It solves the "Two-Language Problem"
 - Come for the speed, stay for the productivity
- You can write high-level code which produces fast, low-level machine code that has traditionally only been generated by static languages.

Benchmarks




More than "Fast R" or "Fast Python"

- Julia is fast because of features which work well together
- You can't just take the magic dust that makes Julia fast and apply it to your favorite language

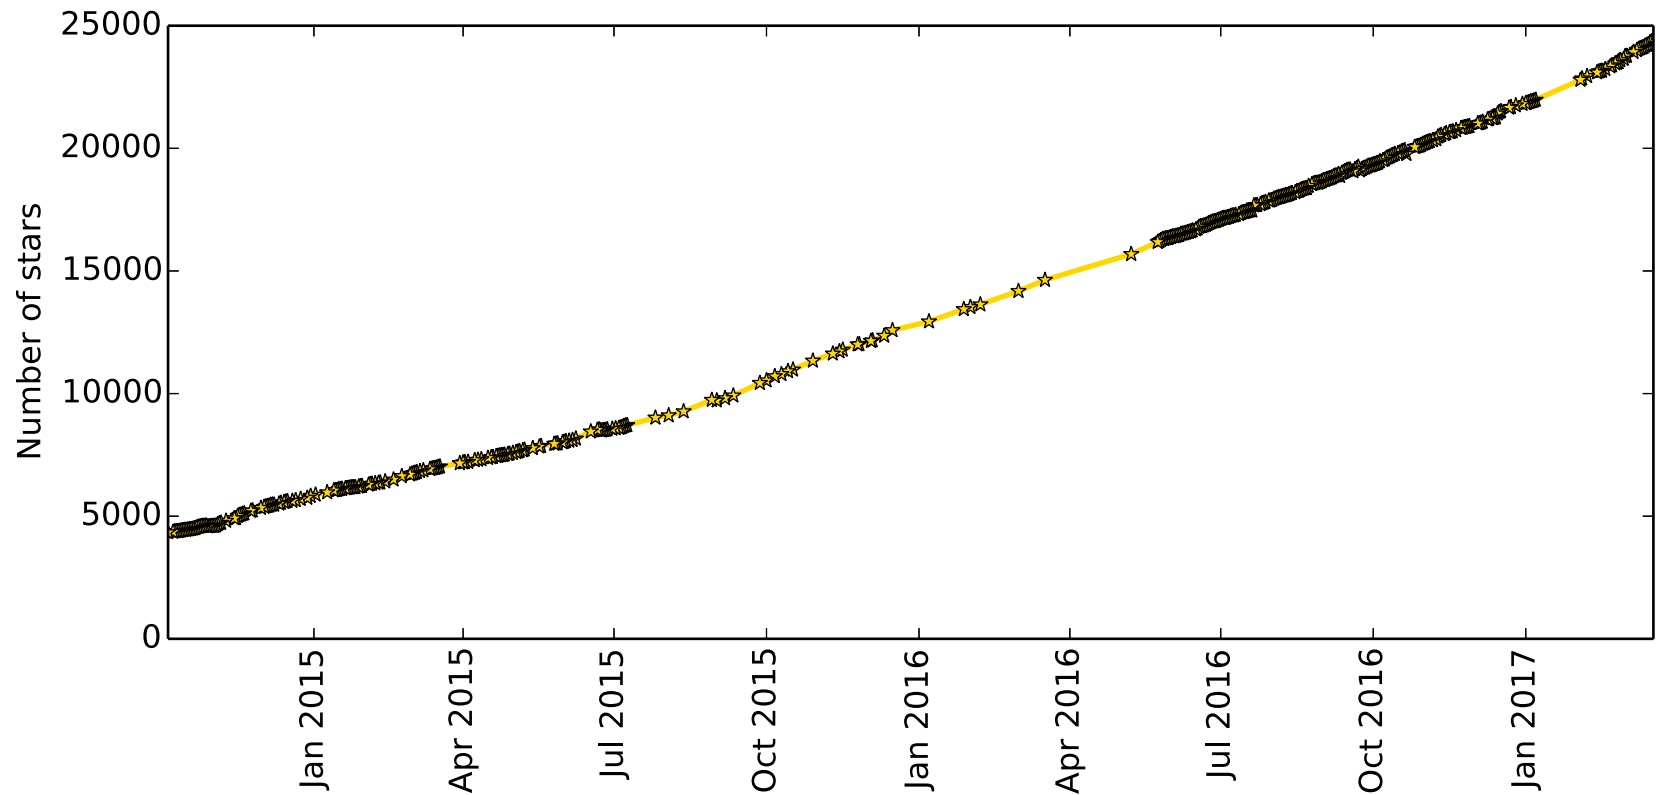
Julia Language Design

- Type system
- Multiple dispatch
- Type inference
- Metaprogramming (macros)
- Just-in-time (JIT) compilation using LLVM
- Clean, familiar syntax
- Function arguments pass by reference

Julia's Growth (Number of Packages)



Julia's Growth (Stars on GitHub)



Design, Type System, and Multiple Dispatch

Julia Basics

Everything has a type

```
1      # typeof(1) == Int64  
  
1.0    # Float64  
  
[1.0, 2.0] # Vector{Float64}
```

Julia Basics

Unicode Characters

- Julia's main focus is numerical computing, so creators wanted code to look similar to mathematical formulas:

$$\hat{\beta} = (X^T X)^{-1} X^T y$$

```
 $\hat{\beta}$  = inv(x'x)x'y  # Implicit multiplication
```

```
 $\hat{\beta}$  = inv(x' * x) * x' * y
```

- Compare this to R:

```
betahat = solve(t(x) %*% x) %*% t(x) %*% y
```

Julia Basics

Functions

```
f(x) = x ^ 2
```

```
# Code blocks require an `end`
```

```
function f(x)
```

```
    x ^ 2
```

```
end
```

Julia Basics

Mutating functions

- Since arguments pass by reference, functions can change arguments in place
- By convention, functions ending in `!` mutate the arguments

```
function makezeros!(v::Array{Float64})  
    v[:] = 0.0  
    v  
end
```


Julia Basics

JIT

- live example with `@code_llvm` macro

```
f(x) = x ^ 2
```

Julia Basics

For loops

Input:

```
for i in 1:3
    println(i)
end
```

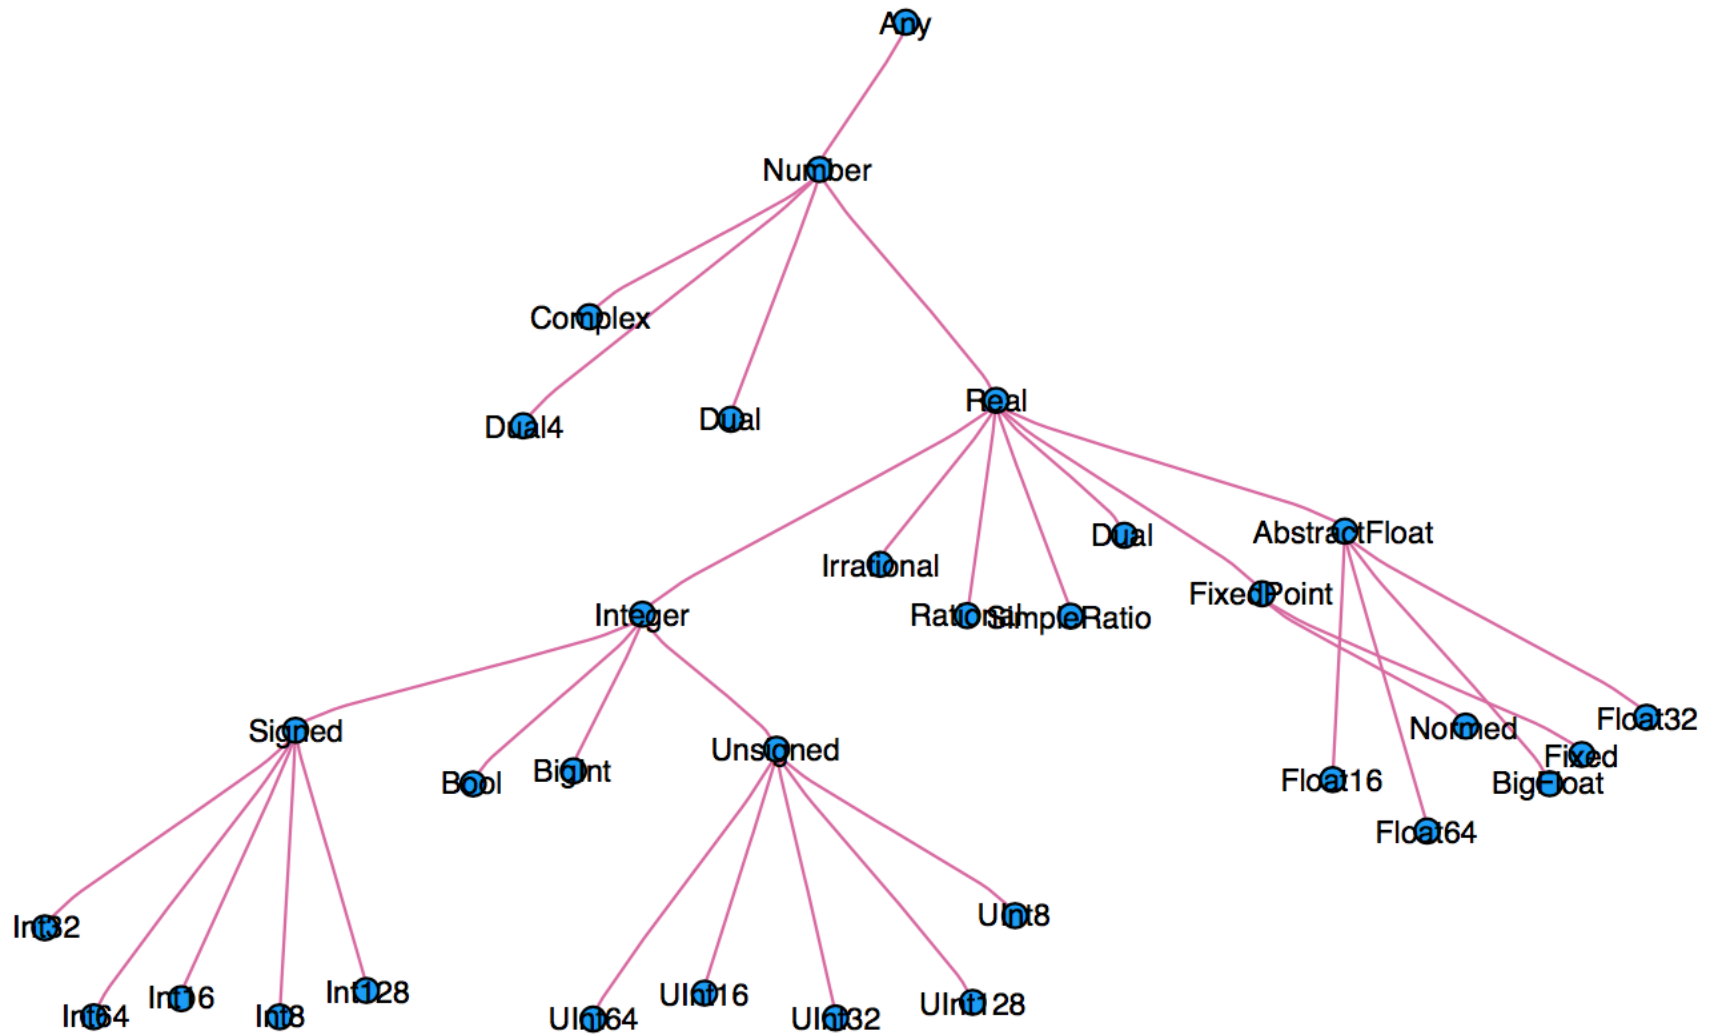
Output:

```
1
2
3
```

Type System

- When thinking about types, think about sets
- An abstract type defines a set of other types
- One abstract type in Julia is `Number`

Number



Number

- What should a `Number` be able to do?

`+`, `-`, `*`, etc.

`methods(*)`

```
# 178 methods for generic function "*":
*(x::Bool, z::Complex{Bool}) in Base at complex.jl:225
*(x::Bool, y::Bool) in Base at bool.jl:91
*(x::Bool, y::T) where T<:Unsigned in Base at bool.jl:104
*(x::Bool, z::Complex) in Base at complex.jl:232
*(x::Bool, y::Irrational) in Base at irrationals.jl:105
*(x::Bool, y::T) where T<:Number in Base at bool.jl:101
*(a::Float16, b::Float16) in Base at float.jl:368
*(x::Float32, y::Float32) in Base at float.jl:374
*(x::Float64, y::Float64) in Base at float.jl:375
...
```

Multiple Dispatch

- The idea of calling different code depending on the types of the arguments is called multiple dispatch.
- Consider the Distributions package, where every distribution has its own type.

```
using Distributions
```

```
mean(Normal(0,1)) == 0.0
```

```
mean(Gamma(10,6)) == 60.0
```

Abstraction

- Consider these three function definitions:

```
# Too broad: Not everything can be added  
f(x) = x + x
```

```
# Too specific: Numbers besides Float64 can be added  
g(x::Float64) = x + x
```

```
# Just right: Every number has a `+` method, so this  
# works on the entire type tree in the previous slide  
h(x::Number) = x + x
```

Quantile Example

- Suppose I want to find quantiles using Newton's method:

$$\theta_{t+1} = \theta_t - \frac{F(\theta_t) - q}{F'(\theta_t)}$$

where F is the CDF of a distribution

- In R, I would need a different function for every distribution!
- In Julia, we can do this in one function

The Power of Julia: Abstraction

- Define functions for the "highest" type you can
- A `UnivariateDistribution` has methods `mean`, `cdf`, `pdf`

```
using Distributions

function myquantile(d::UnivariateDistribution, q::Number)
    θ = mean(d)
    tol = Inf
    while tol > 1e-5
        θold = θ
        θ = θ - (cdf(d, θ) - q) / pdf(d, θ)
        tol = abs(θold - θ)
    end
    θ
end
```

Input:

```
for d in [Normal(), Gamma(5,1), TDist(4)]
    println("For $d")
    println("  > myquantile: $(myquantile(d, .4))")
    println("  > quantile:   $(quantile(d, .4))\n")
end
```

Output:

```
For Distributions.Normal{Float64}(μ=0.0, σ=1.0)
  > myquantile: -0.2533471031356957
  > quantile:   -0.2533471031357997

For Distributions.Gamma{Float64}(α=5.0, θ=1.0)
  > myquantile: 4.1477358804705435
  > quantile:   4.1477358804705435

For Distributions.TDist{Float64}(v=4.0)
  > myquantile: -0.27072229470638115
  > quantile:   -0.27072229470759746
```

Defining types

- Your own types are just as performant as Julia's built-ins.
- Most of Julia is written in Julia

```
# Can change `a` after creating a new instance  
type MyType  
    a::Int64  
end  
  
# Can NOT change `a` after creation  
immutable MyOtherType{T <: Number}  
    a::T  
end
```

Defining methods

Input:

```
f(o::MyType) = o.a + o.a
```

```
o = MyType(4)
```

```
f(o)
```

Output:

```
8
```

Julia's Package Ecosystem

Plotting and Graphics

Julia does not have a built-in plotting package


- [Plots](#)
 - Defines a plotting API that can use several "backends"
- [Gadfly](#)
 - Grammar of graphics for Julia
- Others: PyPlot, UnicodePlots, GR, GLPlot, Winston, ...

Plots with PyPlot

using Plots


`pyplot()` *# use PyPlot backend*

`plot(randn(50, 2))`



Plots with GR

```
gr() # use GR backend  
plot(randn(50, 2))
```



Easy Animations

```
anim = @animate for i in 1:20
    plot(sin, 0, i, xlim = (0,20), ylim = (-1, 1))
end
gif(anim, "/Users/joshday/Desktop/my_animation.gif")
```

INFO: Saved animation to /Users/joshday/Desktop/my_animat.

Calling R from Julia

- Side note: REPL modes
 - `;` (shell mode)
 - `?` (help mode)
- RCall adds
 - `$` (R mode)

```
using RCall  
R"rnorm(5)"
```

Calling Python from Julia

Input:

```
using PyCall  
@pyimport math  
  
math.pi
```

Output:

```
3.141592653589793
```

Statistics, Working with Data, Machine Learning, etc.

- StatsBase, GLM, DataFrames, Query, MixedModels, Distributions, KernelDensity, OnlineStats, LossFunctions,

StatsBase

Much of the functionality built into R

```
using StatsBase
```

```
sample(1:5, 5, replace = false)
```

DataFrames and DataTables

- Both for working with tabular data
- DataTables is a fork of DataFrames
 - Behind the scenes containers are different
 - `dataArray` vs. `NullableArray`

DataTables

Input:

```
iris = readtable(joinpath(Pkg.dir("DataTables"),  
    "test/data/iris.csv"))  
head(iris)
```

Output:

6×5 DataTables.DataTable

Row	SepalLength	SepalWidth	PetalLength	PetalWidth
1	5.1	3.5	1.4	0.2
2	4.9	3.0	1.4	0.2
3	4.7	3.2	1.3	0.2
4	4.6	3.1	1.5	0.2
5	5.0	3.6	1.4	0.2
6	5.4	3.9	1.7	0.4

Query (query almost any data source)

```
using Query
```

```
x = @from i in iris begin  
    @where i.Species == "setosa" && i.PetalLength > 1.7  
    @select i  
    @collect DataTable  
end
```

2×5 DataTables.DataTable

Row	SepalLength	SepalWidth	PetalLength	PetalWidth
1	4.8	3.4	1.9	0.2
2	5.1	3.8	1.9	0.4

Macros

Macros

- Macros are functions of expressions
- They change an expression before it is run, and can therefore do many things

```
x = randn(1000);
```

```
@time sum(x)  # JIT at work  
# 0.031018 seconds (12.86 k allocations: 602.013 KB)  
# -34.195601715147035
```

```
@time sum(x)  
# 0.000003 seconds (5 allocations: 176 bytes)  
# -34.195601715147035
```

Macros (turn off bounds checking)

```
x = rand(1000)

for i in eachindex(x)
    @inbounds x[i] *= 5.0
end
```

Macros (benchmarking)

```
using BenchmarkTools
@benchmark sum(x)
```

```
BenchmarkTools.Trial:
 memory estimate: 16 bytes
 allocs estimate: 1
-----
 minimum time:      223.916 ns (0.00% GC)
 median time:      231.385 ns (0.00% GC)
 mean time:        237.263 ns (0.86% GC)
 maximum time:      7.247 µs (95.52% GC)
-----
 samples:           10000
 evals/sample:      455
 time tolerance:    5.00%
 memory tolerance:  1.00%
```

Macros (Views)

```
x = randn(100,100)
x[:, 1]          # creates a copy
@view x[:, 1]    # creates a "view"
```

Macros (Other)

```
@which sum(x)  # find the method being called
```

```
@edit sum(x)   # open file where the method is
```

Recommendations

My Thoughts

- Julia is ideal for developing projects from scratch/with few dependencies
- The package ecosystem is growing fast but is still lacking much of the functionality you can find in R
- However, there are many interesting things in Julia that R doesn't have:
 - Plots, Convex, JuMP, OnlineStats, Distributions, LossFunctions, ...

Thank You